Mathematical Foundations of Deep Neural Networks, M1407.001200
E. Ryu
Spring 2024

<center>Homework 3
Due 5pm, Monday, March 25, 2024</center>

Problem 1: *3-layer MLP to fit a univariate function.* Consider the univariate function

$$f_\star(x) = (x - 2)\cos(4x).$$

Let $f_\theta(x)$ be a 3-layer MLP with the sigmoid activation function, i.e., $\sigma(x) = (1+e^{-x})^{-1}$. Given data X_i generated as IID unit Gaussians and corresponding labels $Y_i = f_\star(X_i)$ for $i = 1, \ldots, N$, define the loss function

$$\mathcal{L}(\theta) = \frac{1}{N}\sum_{i=1}^{N}\frac{1}{2}(f_\theta(X_i) - Y_i)^2.$$

Use PyTorch to train f_θ with minibatch shuffled cyclic SGD applied to

$$\underset{\theta \in \mathbb{R}^p}{\text{minimize}} \quad \mathcal{L}(\theta).$$

Use layer widths $(n_0, n_1, n_2, n_3) = (1, 64, 64, 1)$, total epochs $K = 1000$ (this is epochs, not iterations), stepsize $\alpha = 0.1$, batch size $B = 128$, and number of data points $N = 512$. (Use the starter code `threelayerSGD.py`.) For all three layers, initialize the weights as IID unit Gaussians and biases with the (deterministic) value 0.03. Plot the final trained function with $f_{\theta_K}(x)$ as a function of x.

Hint. For initialization, do something like

```
model.l1.weight.data = torch.normal(0, 1, model.l1.weight.shape)
model.l1.bias.data = torch.full(model.l1.bias.shape, 0.03)
```

Hint. For the squared loss, use `nn.MSELoss()`.

Solution. See `threelayerSGD_sol.py`. ∎

Problem 2: *Deep learning operates under $p \gg N$.* In the previous problem, how many trainable parameters are in the 3-layer MLP? Repeat the previous problem with the training labels

```
y_train = f_true(X_train) + torch.normal(0, 0.5, X_train.shape)
```

How are the results affected?

Remark. From a classical statistical perspective, it is surprising that large neural networks with more parameters p than the number of data points N do not "overfit", even in the presence of label noise. In fact, most of deep learning operates under the regime where there are more unknowns (trainable parameters) than data points. We will revisit this issue later in our discussion of the bias-variance tradeoff and the double descent phenomenon.

Solution. The total number of trainable parameters is

$$p = (64 + 64) + (64^2 + 64) + (64 + 1) = 4353.$$

This is much more than the number of data points $N = 512$. The noise in training labels do not qualitatively affect the results. ∎

Problem 3: *Basic properties of the CE loss.* Define the *cross entropy loss* as

$$\ell^{\text{CE}}(f, y) = -\log\left(\frac{\exp(f_y)}{\sum_{j=1}^{k} \exp(f_j)}\right),$$

where $f \in \mathbb{R}^k$ and $y \in \{1, \ldots, k\}$.

(a) Show that $0 < \ell^{\text{CE}}(f, y) < \infty$.

(b) Let e_i be the i-th unit vector, i.e, e_i is the one-hot vector with value 1 is the i-th coordinate and 0's for all other coordinates. Show that $\ell^{\text{CE}}(\lambda e_y, y) \to 0$ as $\lambda \to \infty$.

Solution. (a) Since $e^x > 0$ for all $x \in \mathbb{R}$, $0 < \exp(f_y) < \sum_{j=1}^{k} \exp(f_j)$. This gives $0 < \frac{\exp(f_y)}{\sum_{j=1}^{k} \exp(f_j)} < 1$. Thus,

$$0 < \ell^{\text{CE}}(f, y) = -\log\left(\frac{\exp(f_y)}{\sum_{j=1}^{k} \exp(f_j)}\right) < \infty.$$

(b)

$$\ell^{\text{CE}}(\lambda e_y, y) = -\log\left(\frac{\exp(\lambda)}{\exp(\lambda) + k - 1}\right)$$

Since $\lim_{\lambda \to \infty} \frac{\exp(\lambda)}{\exp(\lambda)+k-1} = 1$, $\lim_{\lambda \to \infty} \ell^{\text{CE}}(\lambda e_y, y) = 0$. ∎

Problem 4: *Derivative of max.* Let

$$f(x) = \max\{f_1(x), \ldots, f_N(x)\},$$

where $f_1, \ldots, f_N$ are differentiable univariate functions. Show that if $I = \text{argmax}_{1,\ldots,N}\{f_i(x)\}$ is unique at $x \in \mathbb{R}$ (the maximum is attained by only one function at a given x), then f is differentiable at x and

$$\frac{d}{dx} f(x) = \frac{d}{dx} f_I(x)$$

at x.

Solution. Consider a fixed x_0 such that $I = \text{argmax}_i\{f_i(x_0)\}$ is unique. Since I is unique, $f_I(x_0) > f_i(x_0)$ for all $i \neq I$. Since $f_1, \ldots, f_k$ are differentiable, $f_I - f_i$ is continuous for each i. For each $i \neq I$, there exists an open neighborhood U_i around x_0 such that $f_I(x) > f_i(x)$ for all $x \in x_0$. Since $k < \infty$,

$$U := \bigcap_{i \neq I} U_i$$

is an open neighboorhood on which $f_I(x) > f_i(x)$ for all $x \in x_0$ and $i \neq I$. Thus $f(x) = f_I(x)$ on U, and we can conclude $f'(x_0) = f_I'(x_0)$. ∎

Problem 5: *Basic properties of activation functions.* Prove the following basic facts about some commonly used activation functions.

(a) *Idempotence of ReLU.* The ReLU activation $\sigma(z) = \max\{0, z\}$ is *idempotent*, i.e.,

$$\sigma(\sigma(z)) = \sigma(z), \qquad \forall z \in \mathbb{R}.$$

(b) *Softplus.* The *softplus* function $\sigma(z) = \log(1 + e^z)$ is considered a smooth alternative of ReLU. Show that softplus has Lipschitz continuous derivatives while ReLU does not.

(c) *Equivalence of tanh and sigmoid.* Let $\sigma(z) = (1 + e^{-z})^{-1}$ be the sigmoid function and let $\rho(z) = (1 - e^{-2z})/(1 + e^{-2z})$ be the tanh function. Show that the two activation functions are equivalent in the sense that MLPs built with them are equivalent: given $L > 1$, $A_1, \ldots, A_L$, and $b_1, \ldots, b_L$, there are $C_1, \ldots, C_L$ and $d_1, \ldots d_L$ such that

$y_L = A_L y_{L-1} + b_L$	$y_L = C_L y_{L-1} + d_L$
$y_{L-1} = \sigma(A_{L-1} y_{L-2} + b_{L-1})$	$y_{L-1} = \rho(C_{L-1} y_{L-2} + d_{L-1})$
$\vdots$	$\vdots$
$y_2 = \sigma(A_2 y_1 + b_2)$	$y_2 = \rho(C_2 y_1 + d_2)$
$y_1 = \sigma(A_1 x + b_1),$	$y_1 = \rho(C_1 x + d_1),$

represent identical $x \mapsto y_L$ mappings, and vice versa. Here, $x \in \mathbb{R}^{n_0}$, $A_\ell, C_\ell \in \mathbb{R}^{n_\ell \times n_{\ell-1}}$, $b_\ell, d_\ell \in \mathbb{R}^{n_\ell}$ for $\ell = 1, \ldots, L$, and σ is applied element-wise.

Remark. The "equivalence" of part (c) should not be understood to mean there is no practical difference between the two activation functions. As we will discuss in later in this course, how one initializes neural network parameters is important. When standard initializations are used, tanh is often easier to train compared to sigmoid, due to the fact that the output of tanh is zero-centered.

Solution. (a) $\sigma(z) = \max\{0, z\} = z$ if and only if $z \geq 0$. Since $\sigma(z) = \max\{0, z\} \geq 0$, $\sigma(\sigma(z)) = \sigma(z)$.

(b) $\sigma'(z) = \frac{e^z}{1 + e^z}$ is Lipschitz since

$$\left| \frac{e^{z_1}}{1 + e^{z_1}} - \frac{e^{z_2}}{1 + e^{z_2}} \right| = |z_1 - z_2| \left| \frac{e^{z'}}{(1 + e^{z'})^2} \right| = |z_1 - z_2| \left| \frac{1}{2 + e^{z'} + e^{-z'}} \right| < \frac{1}{4} |z_1 - z_2|$$

for some z' in between z_1, z_2 by the mean value theorem. On the other hand, the derivative of ReLU is not Lipschitz since

$$\sigma'(z) = \begin{cases} 1 & z > 0 \\ 0 & z < 0 \end{cases}$$

and

$$\frac{\sigma'(\varepsilon) - \sigma'(-\varepsilon)}{2\varepsilon}$$

is unbounded as $\varepsilon \to 0$.

(c) The key relation is

$$2\sigma(2z) = 2/(1 + e^{-2z}) = \rho(z) + 1.$$

3

Write $y_1, y_2, ..., y_{L-1}$ for the outputs of $\sigma(z)$ and $y'_1, y'_2, ..., y'_{L-1}$ for the outputs of $\rho(z)$. We will establish the relation $2y_i = y'_i + 1$ for $i = 1, \ldots, L - 1$. First, choose $C_1 = \frac{1}{2}A_1$, $d_1 = \frac{1}{2}b_1$. Then, $2y_1 = y'_1 + 1$. Then, choose $C_i = \frac{1}{4}A_i$, $d_i = \frac{1}{2}b_i + C_i$ for $i = 2, 3, ..., L - 1$. Then by induction,

$$2y_i = 2\sigma(A_i y_{i-1} + b_i) = 2\sigma(4C_i y_{i-1} + 2d_i - 2C_i) = 2\sigma(2C_i y'_{i-1} + 2d_i) = \rho(C_i y'_{i-1} + d_i) + 1 = y'_i + 1.$$

With $2y_{L-1} = y'_{L-1} + 1$, choose $C_L = \frac{1}{2}A_L$, $d_L = b_L + C_L$ to conclude $y_L = y'_L$. ∎

Problem 6: *Vanishing gradients.* Consider the 2-layer neural network

$$f_\theta(x) = u^\mathsf{T}\sigma(ax + b) = \sum_{j=1}^{p} u_j\sigma(a_jx + b_j),$$

where $x \in \mathbb{R}$ and $a, b, u \in \mathbb{R}^p$. Let σ be the ReLU activation function. Using the data $X_1, \ldots, X_N \in \mathbb{R}$ and labels $Y_1, \ldots, Y_N \in \mathcal{Y}$, we train the neural network by solving

$$\underset{\theta \in \mathbb{R}^{3p}}{\text{minimize}} \quad \frac{1}{N}\sum_{i=1}^{N} \ell(f_\theta(X_i), Y_i)$$

with SGD. We assume $\ell(x, y)$ is differentiable in x. Assume the j-th ReLU output is "dead" at initialization in the sense that $a_j^0 X_i + b_j^0 < 0$ for all $i = 1, \ldots, N$. Show that j-th ReLU output remains dead throughout the training.

Remark. The term "vanishing gradients" refers both to the circumstance where the gradient exactly vanishes (as in this problem) and to the circumstance where the gradient becomes extremely small but not zero.

Solution. The key idea is that every SGD update does not move a_j and b_j. Note that

$$\sigma'(z) = \begin{cases} 1 & z > 0 \\ 0 & z < 0 \end{cases}$$

First, it is given that $a_j^0 X_i + b_j^0 < 0$ for all $i = 1, \ldots, N$. Next, given $a_j^k X_i + b_j^k < 0$ for all $i = 1, \ldots, N$, the gradient computation of homework 2 tells us that $\frac{\partial f_\theta(X_i)}{\partial a_j} = 0$, $\frac{\partial f_\theta(X_i)}{\partial b_j} = 0$ for all $i = 1, \ldots, N$. Since SGD performs the updates

$$a_j^{k+1} = a_j^k - \alpha\frac{\partial \ell}{\partial f_\theta}\frac{\partial f_\theta}{\partial a_j} = a_j^k$$

$$b_j^{k+1} = b_j^k - \alpha\frac{\partial \ell}{\partial f_\theta}\frac{\partial f_\theta}{\partial b_j} = b_j^k$$

we conclude $a_j^{k+1} X_i + b_j^{k+1} < 0$ for all $i = 1, \ldots, N$. Thus j-th ReLU output remains dead. ∎

Problem 7: *Leaky ReLU.* The *leaky ReLU* activation function [1] is defined as

$$\sigma(z) = \begin{cases} z & \text{for } z \geq 0 \\ \alpha z & \text{otherwise,} \end{cases}$$

where α is a fixed parameter (α is not trained) often set to $\alpha = 0.01$. Show that leaky ReLU, instead of ReLU, is used in the previous problem, the gradient no longer exactly vanishes.

Solution. Note

$$\sigma'(z) = \begin{cases} 1 & z > 0 \\ \alpha & z < 0, \end{cases}$$

so $\sigma'(z) \neq 0$ for any $z \neq 0$. Remember from homework 2

$$\frac{\partial f_\theta(X_i)}{\partial a_j} = u_j X_i \sigma'(a_j X_i + b_j)$$

$$\frac{\partial f_\theta(X_i)}{\partial b_j} = u_j \sigma'(a_j X_i + b_j)$$

which nonzero vector unless $u_j = 0$. When the weights are randomly initialized, it is unlikely for u_j to remain 0, hence the j-th output will likely not remain dead. ∎

References

[1] A. L. Maas, A. Y. Hannun, and A. Y. Ng, Rectifier nonlinearities improve neural network acoustic models. *ICML Workshop on Deep Learning for Audio, Speech, and Language Processing*, 2013.